

AI Knowledge Assistant

AI Knowledge Assistant.....	1
Problem Statement	1
Overarching Positioning.....	1
Learner Use Case Scenarios	1
1. In-the-Moment Decision Support (High Stakes)	2
2. Just-in-Time Task Execution (Operational Accuracy)	2
3. Skill Reinforcement & Practice (Behavioral Retention)	2
Tech / Experience Considerations	2
Value Proposition.....	4
Right to Win (Strengthened & Differentiated)	4
1. Trusted, Sector-Specific Expertise (25+ Years)	4
2. Deep Organizational Context.....	4
3. Closed-Loop Learning Ecosystem.....	4
4. Compliance, Safety, and Auditability	5
5. Instructionally Designed for Behavior Change.....	5
Strategic Positioning (Clear Contrast with Generic AI)	5
Foundation for Future Transformation	6

Problem Statement

Organizations invest in high-quality, research-backed training to drive critical behaviors, but that training is delivered as a discrete event—disconnected from the real-world moments where those behaviors must be applied.

When those moments arise, learners default to guessing, observing peers, using generic AI tools, or taking no action at all. These alternatives lack contextual relevance, organizational alignment, and verified accuracy—resulting in inconsistent performance, increased risk, and weakened organizational outcomes.

Despite 25+ years of trusted partnerships, tailored content, and validated learning outcomes, there is currently no mechanism to **extend that trusted training into the moment of need**, where it has the greatest impact.

Overarching Positioning

Enable learners to access **real-time, context-aware, organization-aligned guidance** grounded in their assigned training, organizational policies, and role-specific needs—at the exact moment they need to apply a skill.

This is not a generic Q&A tool—it is an **extension of the organization’s training ecosystem into live decision-making environments**.

Learner Use Case Scenarios

1. In-the-Moment Decision Support (High Stakes)

- Student navigating peer pressure around substance use
- Police officer applying de-escalation techniques in the field
- Employee responding to a safety hazard or unsafe behavior

👉 Guidance is:

- Aligned to training they completed
- Consistent with institutional policies and legal standards
- Delivered in a way that is safe, actionable, and defensible

2. Just-in-Time Task Execution (Operational Accuracy) or Resource Access

- Facilities worker recalling steps in a repair process
- Public sector worker following regulated procedures

- Student looking for career readiness or campus support services information

👉 Guidance is:

- Procedurally accurate and role-specific
- Filtered to the exact step or need
- Grounded in approved, auditable content (not internet synthesis)

3. Skill Reinforcement & Practice (Behavioral Retention)


- Practicing difficult conversations
- Reinforcing concepts post-training

👉 Experience includes:

- Scenario-based reinforcement tied to prior training
- Feedback aligned to validated learning objectives
- Continuity between learning, practice, and application

Illustrative Example

We'll have a few sources of data: general knowledge from an LLM, Vector knowledge (primarily courses), and then customer content (a lot of rich types that we have in our different systems). All of those sources are inputs to our magic ai features. The example below shows a customer asking a question about a ladder and how the result may appear for each data source. If we just use an LLM, the answer is pretty generic and we're not adding any real value (aside from being an integrated platform they're accessing for training). If we use our course, we can provide "vector verified" answers based on OSHA. If we use our course and the customers own policy that sits in our system, we get "vector verified" plus their additional context, and the most accurate and valuable answers.




Tech / Experience Considerations

- Don't want learners to have to log in to an LMS to access the Knowledge Assistant – perhaps a mobile app?
 - Considerations: facial scan for quick login, SSO
 - If this was a value add app that customers encouraged their learners to download, it could be used as a mechanism to actually deploy and complete training as well (replacing the need for logging into the LMS)
- The tool will need to know what courses the learner has taken / has access to based on their org's subscription
 - What level of content filtering do we need to support? We should be trying to consume as much of their licensed content as possible to create the broadest knowledge base per client
 - Any course the tenant has access to? (Easiest)
 - Any course the individual learner has assignments for?
 - Including past assignments, or only active?
 - Any course the individual learner has *completed*?
 - What if they've only started it?
- Is this expected to contribute to a learner's assigned learning? This would probably mostly be used as complementary and not a replacement for other assignments, but there may be interest in learners actually being able to complete trainings within this application. It's also possible that certain use cases may want to assign some kind of interaction with the chat tool (ex: field training for new police officers; ongoing training for dispatchers on how they handled calls, etc.) -- need to explore in customer research
 - If they interact with a subject on a topic they've been assigned, but haven't completed is that expected to contribute to their progress in the LMS?

- Is 3rd party content (SCORM, AICC, etc) supposed to feed into this? We should be trying to consume as much of their licensed content as possible
- Is Custom content supposed to feed into this? Ideally, yes. We should also enable customer to load policies, resources, SOPs, or other relevant materials into the tool to train the LLM on their specific instance (their materials cannot be referenced by LLM outside of their instance)
- Ideally content can be served up in a way that isn't just text. For example, if we have a 3D interface of a valve as part of a facilities course, it would be great to be able to extract that in its current state / form within the Knowledge Assistant
- We may want to enable the learner to deep link from the tool into the relevant module of a course based on what they're asking about
- What happens when a prompt goes beyond what we cover in our courses? For something like this, I think we need to be entirely gated by vector content and not the broader world of information the model may have access to.
- Guard rails for high risk situations we can't fully support to help ensure we aren't attempting to guide a user through a situation we can't support.
 - Identify these upfront so user doesn't go down a rabbit hole that doesn't end up being a good experience
- Will need to test responses to prompts (preferably human SMEs but potentially AI agents) to ensure quality
 - From Kenny: We'll need to do some review of the AI answers (AI does it, people do it, we do sample sizes, or something). I suspect when we review, we would assign ratings on the responses, and based on the ratings we will take one of several actions.

Internal review of response (AI or human)



-
- Create capability for Vector or customers to insert soft updates into the LLM (mitigating the need for creating or updating trainings or formal documentation) when learner inputs prompts that are not covered by our knowledge base
- Future Considerations
 - Upsell feature – Usage feature that suggests other content based on customer's specific use cases
 - Admin feature to build company specific flows – similar to end user use case, but admin can enter a scenario, review output and then provide further guidance/context to adjust the output further.
 - Users would have a list of admin created scenarios to select or admin could assign

Value Proposition

Extend trusted, compliant training into real-world moments—delivering context-aware, organization-aligned guidance that drives consistent, high-quality decisions when it matters most.

Right to Win (Strengthened & Differentiated)

Your advantage is no longer just *what you know*, but **how deeply you are embedded in the learning, organizational, and operational context.**

1. Trusted, Sector-Specific Expertise (25+ Years)

- Longstanding partnerships across education, commercial, and public sectors
- Proven credibility as a source of compliant, research-backed content
- Institutional trust that generic AI cannot replicate

2. Deep Organizational Context

- Training is:
 - Assigned to specific learner groups
 - Delivered at intentional moments
 - Customized with organization-specific policies, resources, and priorities

👉 This enables:

- Responses tailored not just to *a user*, but to *their organization, role, and requirements*

3. Closed-Loop Learning Ecosystem

- Courses → Assessments → Completion data → Surveys → Insights

👉 This creates a system where the assistant can:

- Reinforce what was actually learned (not generic knowledge)
- Adapt to gaps in understanding
- Feed usage data back into content and program improvement

4. Compliance, Safety, and Auditability

- Content aligned to legal and regulatory standards
- Guidance consistent with institutional policies
- Ability to provide **defensible, auditable outputs** in high-stakes scenarios

👉 Critical differentiator vs. generic AI in:

- Public sector
- Workplace safety
- Compliance-driven environments

5. Instructionally Designed for Behavior Change

- Not just information retrieval, but:
 - Behavior modeling
 - Scenario-based learning
 - Skill application

👉 Positions the assistant as:

- A **performance support system**, not just an answer engine

Strategic Positioning (Clear Contrast with Generic AI)

Generic AI:

- Broad, fast, and convenient
- Lacks verified accuracy guarantees
- No organizational alignment
- No accountability

Your AI Knowledge Assistant:

Grounded in trusted content, aware of organizational context, aligned to policy, and designed to drive real-world behavior—not just provide answers.

Foundation for Future Transformation

This initiative is not just a feature—it lays groundwork for:

- **Dynamic, personalized learning pathways**
- **Competency-based progression models**
- **Real-time performance analytics tied to outcomes**
- **Deeper system integration into customer workflows**

👉 In other words:

We move from **“training provider”** → **“continuous performance enablement platform.”**